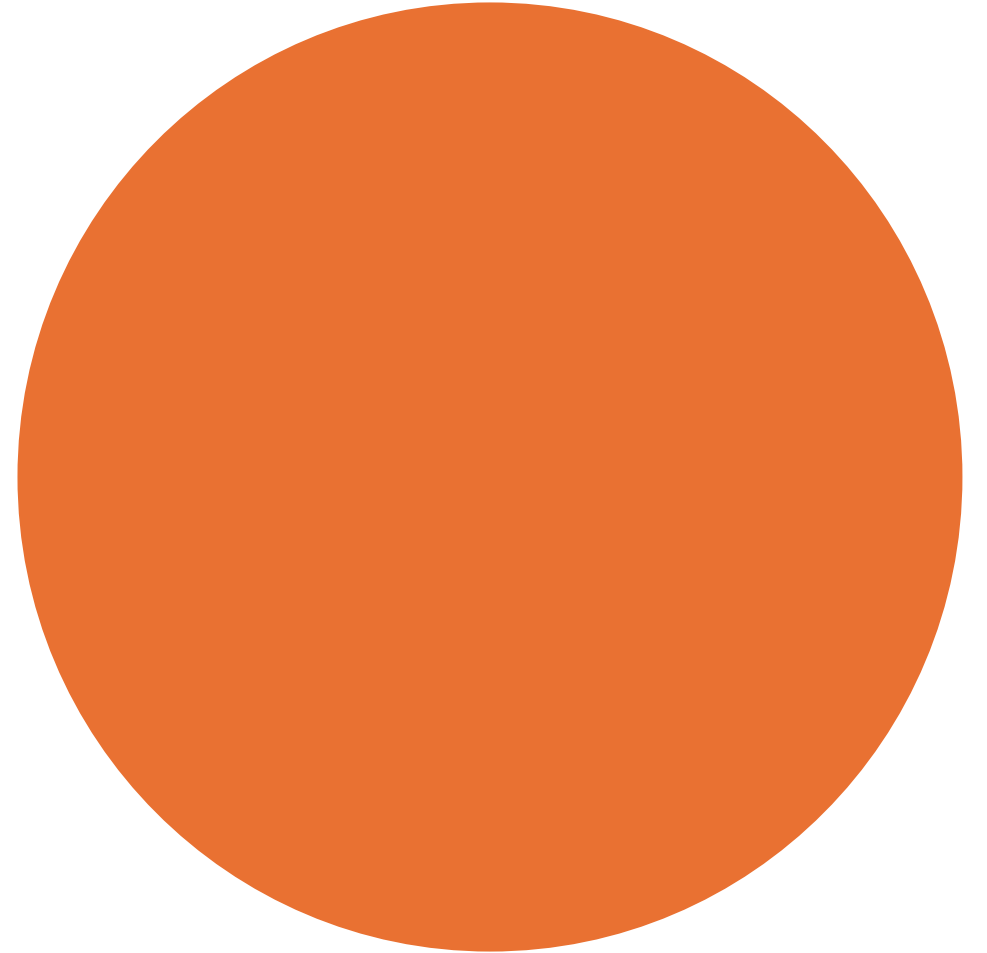


Ch6: Processes



Processes and Programs

A process is an instance of an executing program.

A program is a file containing a range of information that describes how to construct a process at run time.

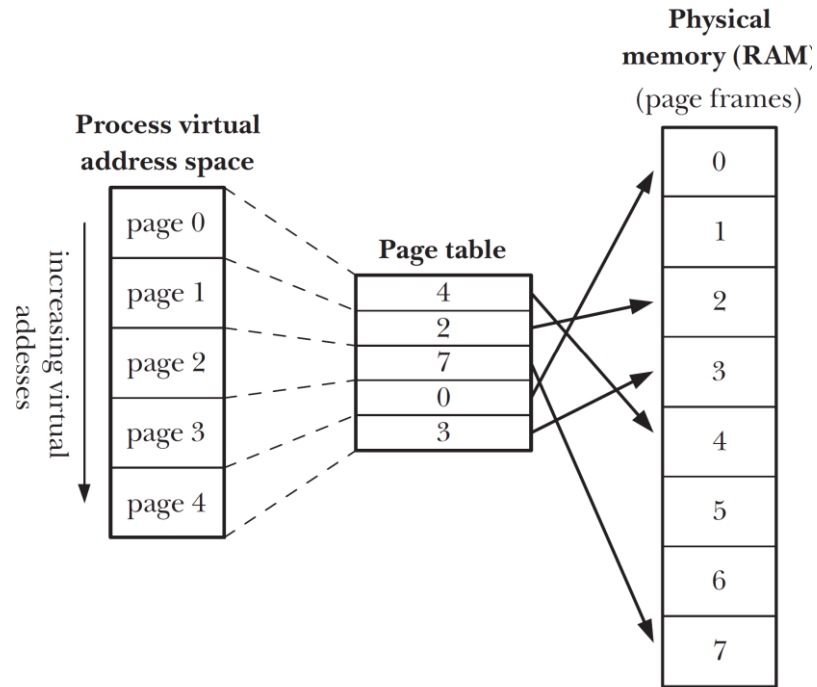
Information exists in a program file:

- Binary format identification (a.out, COFF, ELF).
- Machine-language instructions.
- Program entry-point address.
- Data (initial values and literal strings).
- Symbol and relocation tables.
- Shared-library and dynamic-linking information.

Process components

- **A process consists of:**
 - User-space memory containing program code and variables.
 - kernel data structures that maintain information about the state of the process:
 - Process IDs (/proc/sys/kernel/pid_max, getpid(), getppid()).
 - Virtual memory tables.
 - table of open file descriptors.
 - Resource usage and limits.
 - Others.

Virtual Memory Management



Locality of reference: a typical property of most programs.

- Spatial locality: **reference memory address nearby.**
- Temporal locality: **access same address in near future.**

We can execute a program while only part of its address space in RAM.

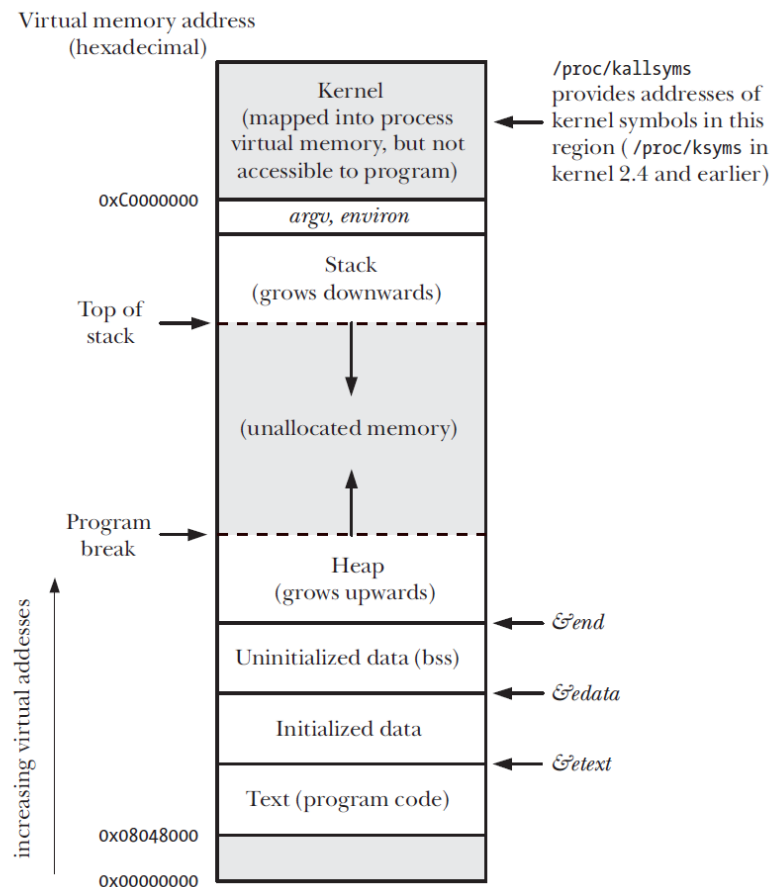
How does it work?

- Memory is divided into pages.
- Kernel maintains a page table for each process.
- Vary over program lifetime (stack, heap, shared memory).
- Page faults.
- Swap area.
- MMU.

Virtual Memory Advantages

- Process isolation (from other process and from the kernel).
- Memory sharing.
 - Executing the same program.
 - IPC.
- Memory Protection (read-only, execute-only, RW).
- Compiler and linker don't need to be concerned with the physical layout.
- Loading programs faster.
- Better CPU utilization.

Memory Layout of a Process



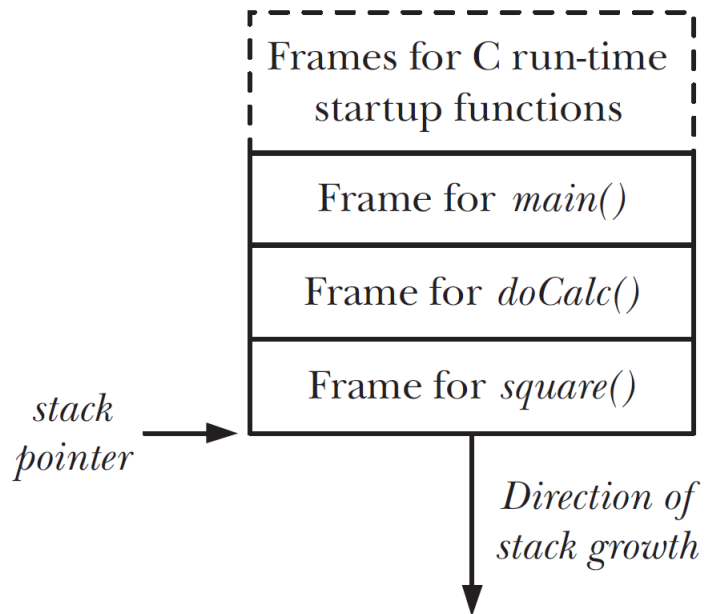
Process memory layout consists of:

- Text (read-only, sharable).
- Data.
- BSS (Block Started by Symbol).
- Stack.
- Heap (Program Break)

Notes:

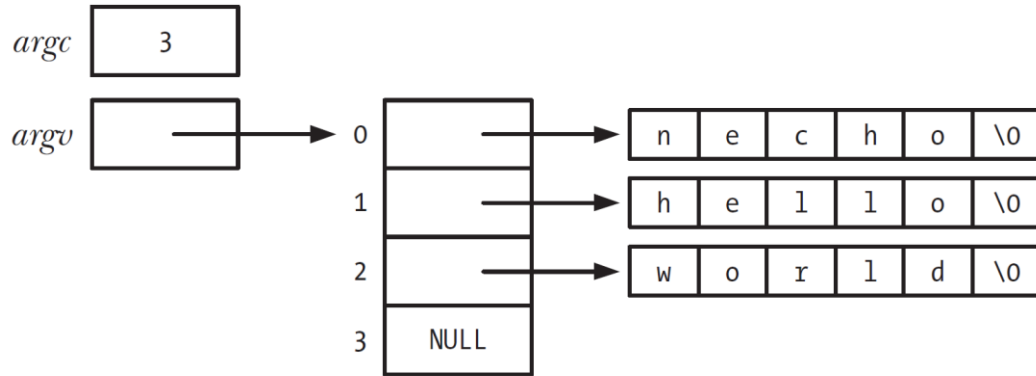
- Size command.
- `etext`, `edata`, `end`.
- Reentrant functions.

Stack and Stack Frames



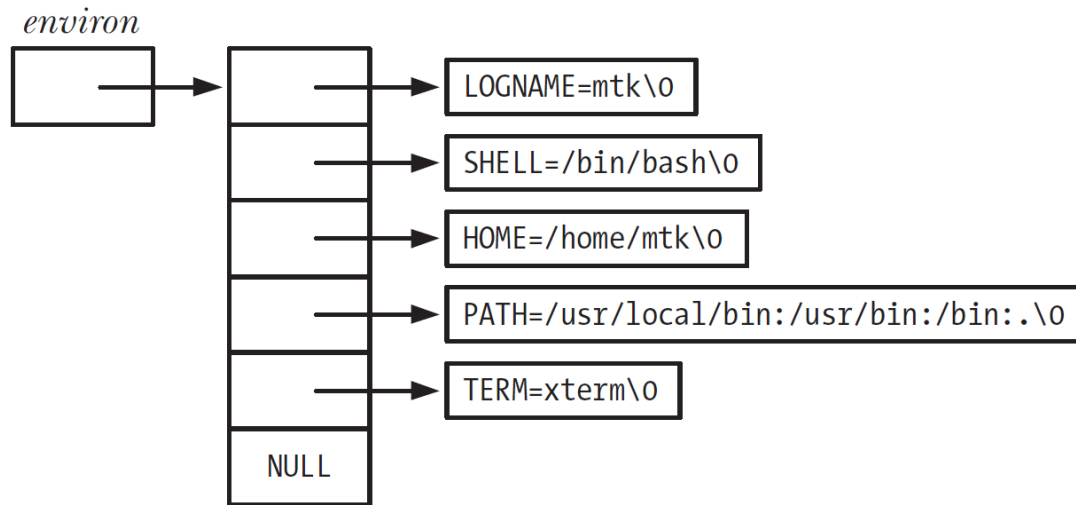
- User Stack vs. Kernel Stack.
- Stack Frames
- Stack Frame contents
 - Function arguments.
 - Local(automatic) variables.
 - Call linkage info.
- Recursion.

Command Line Arguments



- Using argv[0] (gzip, gunzip, zcat).
- /proc/PID/cmdline.
- program_invocation_name.
- program_invocation_short_name.
- Place of argv and environ arrays.
- getopt().

Environment List



- Place of argv and environ arrays.
- One-way, one-time transfer from parent to child.
- environ variable.
- It can be passed to main() but with limited scope.
- /proc/PID/environ.
- export, printenv, unset commands.
- Modifying env in C: setenv, unsetenv.
- Clear env (possible leaks).

Ch7: Memory Allocation



Heap

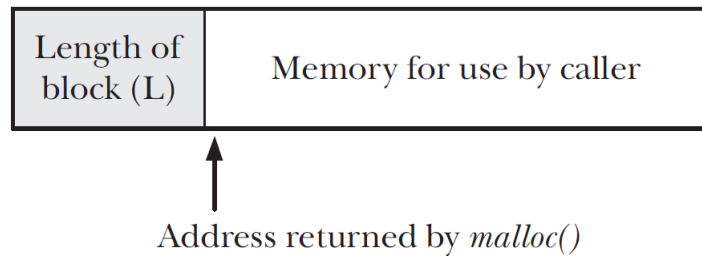
- **The heap** is a variable size segment of contiguous virtual memory that begins just after the uninitialized data segment of a process and grows and shrinks as memory is allocated and freed.
- **Program Break** is The current limit of the heap.
- **Adjusting Program break** using `brk()` and `sbrk()`. Kernel allocates pages when the new memory is accessed.
- Program break upper and lower limits (end of data, shared memory, process limits).
- Getting the current program break (`sbrk(0)`).

Allocating Memory on the Heap: malloc() and free()

- **malloc()** allocates the needed size and may adjust the program break.
- **free()** usually will not lower the program break.
 - Freed memory might be in the middle of the heap.
 - Minimize the number of sbrk() calls.
 - Programs tend to repeatedly release and reallocate memory.
- When a process terminates, all of its memory is returned to the system. Should we call free()?
 - Readability and maintainability.
 - Long-running programs (Daemons and shells).


Implementation of malloc()

- Search the free blocks for a good candidate (first-fit, best-fit, etc.).
 - Split the block if its size is larger than needed.
- Call sbrk() if no free block matching (with larger size).
- How free() will know the size of the block?




Implementation of free()

- Get the block length allocated and set by malloc().
- Construct the free blocks list.



- The free list will become intermingled with blocks of allocated, in-use memory:



Block on free list:

L	P	N	
---	---	---	--

Allocated, in-use block:

L	
---	--

"-" = pointer value marking end of list

Calloc() and realloc()

- **calloc()** allocates memory for an array of identical items and initializes the memory to zero.
- **realloc()** resizes a block of allocated memory.
 - Attempts to coalesce the block with an immediately following block of memory.
 - Might copy all existing data from old to new block.
 - Must use the returned pointer.
 - Do not assign directly as realloc might fail.

Surviving Rules in Dynamic memory Allocation

- DO NOT touch any memory outside the allocated block range.
- DO NOT free an allocated block twice.
- Free with the same pointer returned from malloc. NOT with offset.
- Free the allocated memory.

Allocating Memory on the Stack

→ `alloca()`

- Allocates memory dynamically on the stack.
- Obtains memory from the stack by increasing the size of the stack frame.
- DO NOT call `free()`.
- Allocated memory is valid only within the function.

Debugging Dynamic Memory Allocation

- **mtrace(), muntrace(), MALLOC_TRACE**
 - record tracing information about memory allocation and deallocation.
- **mcheck(), mprobe()**
 - Perform consistency checks on blocks of allocated memory.
- **malloc debugging libraries**
 - Valgrind
 - Electric Fence
 - Insure++
 - dmalloc